



Artificial Intelligence ✓

Opinion +

Technology +

Enough.




Pablo Fuchsberger · 6 min read · May 15, 2026



5



Pablo Fuchsberger, May 2026



In the middle of the most expensive race in tech history, nobody stops to ask whether we've already arrived. Or worse — whether arriving even matters.

There's a question nobody in Silicon Valley wants to say out loud. Asking it means admitting the game might end sooner than expected — and with it, the stratospheric valuations, the cosmic vision speeches, and the photos of CEOs posing in front of server racks like generals in a holy war.

The question is simple: what if we already have enough?

While data centers spread like silicon tumors across the Nevada desert, while tens of billions burn on chips that consume the energy of entire cities, while engineers chase improvements of 30%, 40%, or “infinitely better” — nobody stops the machine to ask whether the destination justifies the journey. Or whether the destination exists at all.

“The race toward AGI looks more and more like a race toward the horizon: the faster you run, the farther it seems. And the more expensive the ticket.”

The goal nobody can define

Artificial General Intelligence — AGI — is the official objective. Imposing. Inevitable. Transformative. The problem is that if you ask ten people to define it precisely, you get ten different answers. For some, it's “a machine that can learn any human cognitive task.” For others, “true causal reasoning.” For others, “consciousness.” For Sam Altman, apparently, it's whatever OpenAI decides it is at whatever moment happens to be most convenient.

That's the core problem: AGI is not a clear technical threshold. It's a philosophical concept dressed up as an engineering objective. And that allows any lab to proclaim it whenever it suits them, and then sell “AGI 2.0” with persistent memory, simulated emotions, spatial reasoning, or whatever else the market needs to hear that quarter. The race has no finish line because the finish line is drawn by whoever is running. And they charge by the mile.

Yann LeCun has been saying it for years, with a bluntness that irritates half the field: language models are not going to reach AGI no matter how much we scale them. Something deeper is missing. A world model. Causal understanding. Real planning under uncertainty. His bet on Joint Embedding Predictive Architectures goes in a completely different direction. He might be right. He might not be. But at least someone is asking the hard questions instead of announcing milestones at conferences with dramatic background music.

What we've already built is, objectively, extraordinary

Let's pause for an uncomfortable moment. In the frenzy over what's missing, we forget how quickly what already exists has blown past everything we expected just three years ago.

Current models reason, synthesize, code, translate, diagnose, and create at a level that surpasses the vast majority of humans in almost any measurable intellectual task. We're not talking about chess or Go anymore — that was solved a decade ago and got maybe one board dinner's worth of attention. We're talking about complex legal reasoning, differential medical diagnosis, creative writing, and university-level mathematics. A modern LLM passes the bar exam, medical board certification, and MIT admissions. Without ever sleeping. Without getting emotional. Without asking for a raise.

Einstein died in 1955. If he could sit down today and discuss general relativity with a language model, he'd find a peer. That's not a minor achievement. It's almost unbelievable. And we treat it like a beta version of something better.

“In eighteen months, we went from the most closely guarded vault in tech to downloading compressed human knowledge onto your laptop, free, on a Sunday afternoon.”

Human knowledge, on your computer, for free

Two years ago, language models were corporate state secrets: trained on private clusters, served behind paid APIs, protected by multi-layered legal walls. A model's weights were the most guarded treasure in the industry. Leaking them was practically an act of piracy.

Today you can download Llama, Mistral, Qwen, or Gemma onto a decent consumer GPU and run them locally. No internet. No sending your data anywhere. No monthly subscription. The synthesized knowledge of all of human civilization, available in your home, 24 hours a day, without asking anyone's permission.

This collapse — from the locked vault to the free download — is one of the fastest movements in the history of any technology. And it carries a consequence the industry hasn't fully processed yet: LLMs are already a commodity. Now. Not in five years. Now.

For 99.9% of real-world use cases, the difference between a model that reasons at PhD level and one at postdoc level is completely irrelevant. Most people don't need an assistant that wins the Nobel Prize. They need help

writing a difficult email, understanding a contract clause, debugging code, or preparing for a meeting. For that, what exists today is already extraordinarily more than enough. And it costs less than a coffee at any airport in the world.

Who's paying the bill? And why?

Training frontier models costs between \$100 million and \$500 million per run. The infrastructure to serve them at global scale adds billions more. Investment plans for the coming years talk about trillions — with a T — in data centers, chips, and energy. Trillions. The kind of number that stops being money and starts being ideology.

Against that, it's worth asking without anesthesia: what business model justifies this investment? Consumer subscriptions are too small. Enterprise APIs are growing, but not at that speed. Government contracts are slow and political. And open source — the most honest part of the ecosystem — generates no direct revenue for whoever publishes.

The implicit bet is that when AGI arrives, it will create so much economic value that everything will have been worth it. But if AGI is undefinable, we're betting on an asset nobody can seriously value. It's like investing trillions to find the end of the rainbow because someone promised gold. And if what we already have is enough to transform the global economy — which it is, and then some — the race stops being technological ambition and starts looking more like an addiction with very good public relations.

This isn't an argument against innovation. It's one against excess.

Research into new architectures, causal reasoning, computational efficiency, alignment, and safety deserves to continue. Without question. Science doesn't owe anyone short-term profitability, and history is full of advances that seemed absurd until they weren't.

But there is a vast difference between genuine research and a corporate arms race disguised as technological philanthropy. There is a difference between pushing the boundaries of knowledge and burning planetary resources — energy, water, minerals, human time — to win a race whose finish line nobody can define because the people organizing it reserve the right to move it.

The industry has a scale problem, not a direction problem. The direction might be right. The scale is a bet on delirium.

“The real risk isn’t that AI goes too far. It’s that we’ve gone so far, spending so much, that we can no longer afford to stop and look at what we built.”

We already arrived. Nobody celebrated.

Three years ago, having access to a system that could reason, write, analyze, and create at the level of a competent professional would have seemed like science fiction. Today you download it on a Sunday. You use it for an email. And by Monday they’re already announcing the next model is 40% better on benchmarks nobody uses in real life.

The problem isn’t ambition. It’s that ambition became the end in itself. That stopping, looking at what was built, and saying “this is already enough to change the world — now let’s use it well” has become something nobody in the industry can say without looking like a coward or a failure.

But sometimes the hardest act isn’t to keep running. It’s to have the honesty — and the maturity — to stop.

We already arrived. What we do with that is the only question that matters.

. . .

A note on how this article was written: the ideas, arguments, and opinions here are entirely my own — born from years of following this industry closely and, frankly, from growing frustration. But I used an LLM to help sharpen the prose, tighten the structure, and push the tone further than I might have on my own. I find it only fair to say so openly, especially in a piece that argues we should take seriously what these tools can already do. If the writing is better than usual, that’s partly the point. If the argument holds up, that’s on me.

Artificial Intelligence



Opinion



Technology



Written by Pablo Fuchsberger

22 followers · 80 following

Follow



Published in DataDrivenInvestor

110K followers · Last published 1 day ago

Following

empowerment through data, knowledge, and expertise. Join DDI community at <https://join.datadriveninvestor.com>